

Problem 1

1.

```
(xray) jason (master) pytorch-ssd $ python eval_ssd.py --net mb1-ssd --dataset ./data/VOC2007test/ --trained_model models/mobilenet-v1-ssd-mp-0.67  
$ pth --label_file models/voc-model-labels.txt
```

```
process image 4950  
Load Image: 0.003893 seconds.  
Inference time: 0.04917764663696289  
Prediction: 0.067090 seconds.  
process image 4951  
Load Image: 0.004049 seconds.  
Inference time: 0.048860788345336914  
Prediction: 0.070375 seconds.  
  
Average Precision Per-class:  
aeroplane: 0.6742489426027927  
bicycle: 0.7913672875238116  
bird: 0.612096015101108  
boat: 0.5616407126931772  
bottle: 0.3471259064860268  
bus: 0.7742298893362103  
car: 0.7284171192326804  
cat: 0.8360675520354323  
chair: 0.5142295855384792  
cow: 0.6244090341627014  
diningtable: 0.7060025454924147  
dog: 0.7849252606216821  
horse: 0.8202146617282785  
motorbike: 0.793578272243471  
person: 0.7042670984734087  
pottedplant: 0.40257147509774405  
sheep: 0.6071252282334352  
sofa: 0.7549120254763918  
train: 0.8270992920206008  
tvmonitor: 0.6459903029666852  
  
Average Precision Across All Classes: 0.6755259103533267
```

2.

```

Inference time: 0.008140802383422852
Prediction: 0.045243 seconds.
process image 120
Load Image: 0.032561 seconds.
Inference time: 0.008483171463012695
Prediction: 0.049908 seconds.
process image 121
Load Image: 0.036689 seconds.
Inference time: 0.008579492568969727
Prediction: 0.048602 seconds.
process image 122
Load Image: 0.018957 seconds.
Inference time: 0.008281946182250977
Prediction: 0.051239 seconds.

Average Precision Per-class:
Handgun: 0.015981481468796188
Shotgun: 0.0051717401545078754

Average Precision Across All Classes:0.010576610811652032

```

no finetuning

```

Inference time: 0.11897158622741699
Prediction: 0.123757 seconds.

Average Precision Per-class:
Handgun: 0.8270707717875698
Shotgun: 0.5643195135016551

Average Precision Across All Classes:0.6956951426
446125

```

sorry it is hard to see pretrained handgun is 0.82 shotgun is 0.56

3.

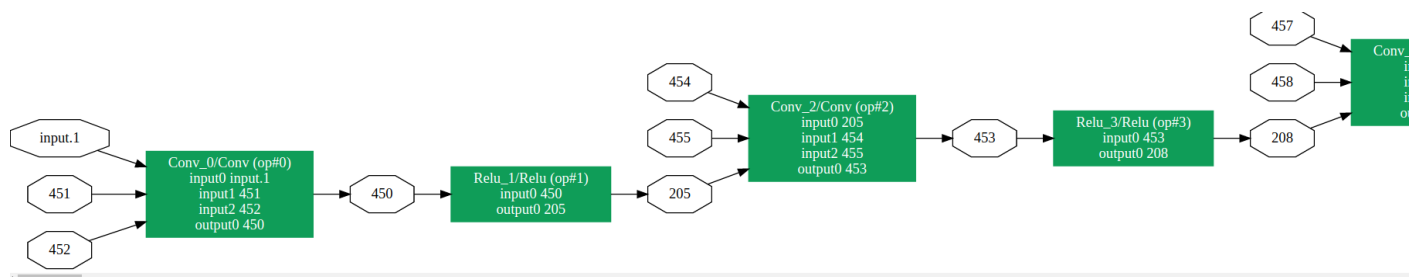
```

(testssd) jason (master *) pytorch-ssd $ python convert_to_caffe2_models.py mb1-ssd mb1-ssd-Epoch-95
-Loss-2.7047196984291078.pth models/open-images-model-labels.txt

```

[onnx file \(https://drive.google.com/file/d/1hDzMHGq2f3GzQmGMdvr9mbuROuwqr3w1/view?usp=sharing\)](https://drive.google.com/file/d/1hDzMHGq2f3GzQmGMdvr9mbuROuwqr3w1/view?usp=sharing)

4.



[link to svg \(https://drive.google.com/file/d/1crto_mLPnFmmilJaels50hoENDgGhzul/view?usp=sharing\)](https://drive.google.com/file/d/1crto_mLPnFmmilJaels50hoENDgGhzul/view?usp=sharing)

5.

```

(testssd) jason projects $ sudo docker run -it -v $(pwd):$(pwd) -p 127.0.0.1:9001:8001 mcr.microsoft
com/onnxruntime/server --model_path $(pwd)/mb1-ssd.onnx

```

```
In [42]: # Import some dependency libraries that we are going to need to run the

import numpy as np
import requests
from PIL import Image
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import onnx_ml_pb2 as onnx_ml_pb2
import predict_pb2 as predict_pb2
```

```
In [43]: # Load the raw image

input_shape = (1, 3, 300, 300)
img = Image.open("handgun.jpg")
img = img.resize((300, 300), Image.BILINEAR)

# Let us see what the input image looks like
img
```

Out[43]:



```
In [44]: img_data = np.array(img)
img_data = np.transpose(img_data, [2, 0, 1])
img_data = np.expand_dims(img_data, 0)
mean_vec = np.array([0.485, 0.456, 0.406])
stddev_vec = np.array([0.229, 0.224, 0.225])
norm_img_data = np.zeros(img_data.shape).astype('float32')
for i in range(img_data.shape[1]):
    norm_img_data[:,i,:,:] = (img_data[:,i,:,:]/255 - mean_vec[i]) / st
```

```
In [45]: # Create request message to be sent to the ORT server

input_tensor = onnx_ml_pb2.TensorProto()
input_tensor.dims.extend(norm_img_data.shape)
input_tensor.data_type = 1
input_tensor.raw_data = norm_img_data.tobytes()

request_message = predict_pb2.PredictRequest()

# For your model, the inputs name should be something else customized by
request_message.inputs["input.1"].data_type = input_tensor.data_type
request_message.inputs["input.1"].dims.extend(input_tensor.dims)
request_message.inputs["input.1"].raw_data = input_tensor.raw_data

content_type_headers = ['application/x-protobuf', 'application/octet-st

for h in content_type_headers:
    request_headers = {
        'Content-Type': h,
        'Accept': 'application/x-protobuf'
    }
```

```
In [46]: PORT_NUMBER = 9001 # Change appropriately if needed based on any change.
inference_url = "http://127.0.0.1:" + str(PORT_NUMBER) + "/v1/models/de
response = requests.post(inference_url, headers=request_headers, data=r
```

```
In [46]: PORT_NUMBER = 9001 # Change appropriately if needed based on any change.
inference_url = "http://127.0.0.1:" + str(PORT_NUMBER) + "/v1/models/de
response = requests.post(inference_url, headers=request_headers, data=r
```

```
In [47]: # Parse response message

response_message = predict_pb2.PredictResponse()
response_message.ParseFromString(response.content)

# For your model, the outputs names should be something else customized
boxes = np.frombuffer(response_message.outputs['boxes'].raw_data, dtype=
#labels = np.frombuffer(response_message.outputs['labels'].raw_data, dt
scores = np.frombuffer(response_message.outputs['scores'].raw_data, dt

print('Boxes shape:', response_message.outputs['boxes'].dims)
#print('Labels shape:', response_message.outputs['labels'].dims)
print('Scores shape:', response_message.outputs['scores'].dims)

Boxes shape: [1, 3000, 4]
Scores shape: [1, 3000, 3]
```

```

response_message.ParseFromString(response.content)

# For your model, the outputs names should be something else customized
boxes = np.frombuffer(response_message.outputs['boxes'].raw_data, dtype=
#labels = np.frombuffer(response_message.outputs['labels'].raw_data, dtype=
scores = np.frombuffer(response_message.outputs['scores'].raw_data, dtype=

print('Boxes shape:', response_message.outputs['boxes'].dims)
#print('Labels shape:', response_message.outputs['labels'].dims)
print('Scores shape:', response_message.outputs['scores'].dims)

```

```

Boxes shape: [1, 3000, 4]
Scores shape: [1, 3000, 3]

```

```

In [54]: ## Display image with bounding boxes and appropriate class

# Parse the list of class labels
classes = [line.rstrip('\n') for line in open('labels.txt')]

# Plot the bounding boxes on the image
plt.figure()
fig, ax = plt.subplots(1, figsize=(12,9))
ax.imshow(img)

resized_width = 300 # we resized the original image, remember ?
resized_height = 300
num_boxes = 6 # we limit displaying to just 10 boxes to avoid clogging
                # The results are already sorted based on box confidence.

for c in range(num_boxes):
    base_index = c * 4
    y1, x1, y2, x2 = boxes[base_index] * resized_height, boxes[base_index] * resized_width
    color = 'blue'
    box_h = (y2 - y1)
    box_w = (x2 - x1)
    bbox = patches.Rectangle((y1, x1), box_h, box_w, linewidth=2, edgecolor=color)
    ax.add_patch(bbox)
    #plt.text(y1, x1, s=classes[labels[c] - 1], color='white', verticalalignment='top')
plt.axis('off')

# Save image
#plt.savefig("output/ssd_result.jpg", bbox_inches='tight', pad_inches=0)
plt.show()

```

<Figure size 432x288 with 0 Axes>

6. Couldnt figure out what was wrong with my code for finding the

```
for c in range(num_boxes):  
    base_index = c * 4  
    y1, x1, y2, x2 = boxes[base_index] * resized_height, boxes[base_index] * resized_width  
    color = 'blue'  
    box_h = (y2 - y1)  
    box_w = (x2 - x1)  
    bbox = patches.Rectangle((y1, x1), box_h, box_w, linewidth=2, edgecolor=color)  
    ax.add_patch(bbox)  
    #plt.text(y1, x1, s=classes[labels[c] - 1], color='white', verticalalignment='top',  
plt.axis('off')  
  
# Save image  
#plt.savefig("output/ssd_result.jpg", bbox_inches='tight', pad_inches=0)  
plt.show()
```

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



Problem 2

1.

GCP - vertex ai, i am including the deep learning frameworks as it is included in the listed [link \(https://cloud.google.com/deep-learning-vm/docs/images\)](https://cloud.google.com/deep-learning-vm/docs/images)

- tensorflow enterprise 2.x, 1.x
- pytorch - latest version
- pytorch XLA - latest
- Chainer - latest
- MXNet - latest
- CNTK - latest
- Caffe - latest

Microsoft - azure datascience vm [link \(https://docs.microsoft.com/en-us/azure/machine-learning/data-science-virtual-machine/dsvm-tools-deep-learning-frameworks\)](https://docs.microsoft.com/en-us/azure/machine-learning/data-science-virtual-machine/dsvm-tools-deep-learning-frameworks) azure ml [link \(https://docs.microsoft.com/en-us/azure/machine-learning/overview-what-is-azure-machine-learning\)](https://docs.microsoft.com/en-us/azure/machine-learning/overview-what-is-azure-machine-learning)

- pytorch - version 1.9.0
- tensorflow - version 2.5

IBM - cloud pak - [link \(https://www.ibm.com/docs/en/cloud-paks/cp-data/3.0.1?topic=models-supported-frameworks\)](https://www.ibm.com/docs/en/cloud-paks/cp-data/3.0.1?topic=models-supported-frameworks)

- XGBOOST - 0.8, 0.9
- Tensorflow - 1.15, 2.1
- Caffe - 1.0
- Keras 2.2.5
- Pytorch, 1.1, 1.2, 1.3

Amazon - deeplearning containers [link \(https://github.com/aws/deep-learning-containers/blob/master/available_images.md\)](https://github.com/aws/deep-learning-containers/blob/master/available_images.md)

- tensorflow
- pytorch 1.10.0, 1.9.0, etc...
- tensorflow - 2.7.0, 2.6.2, 2.6.0
- chainer - 4.0.0, 4.1.0, 5.0.0
- mxnet - 1.8.0, 1.7.0, etc..

2.

Google Vertex AI

- NVIDIA_TESLA_A100
- NVIDIA_TESLA_K80
- NVIDIA_TESLA_P4
- NVIDIA_TESLA_P100
- NVIDIA_TESLA_T4
- NVIDIA_TESLA_V100

Amazon

- NVIDIA_TESLA_A100
- NVIDIA_TESLA_M60
- NVIDIA_T4
- NVIDIA_TESLA_V100
- NVIDIA_A10G

IBM [link \(https://www.ibm.com/downloads/cas/RDPDBJ3X\)](https://www.ibm.com/downloads/cas/RDPDBJ3X).

- NVIDIA_TESLA_M60
- NVIDIA_TESLA_K80
- NVIDIA_TESLA_P100
- NVIDIA_TESLA_V100

Microsoft Azure [link \(https://azure.microsoft.com/en-us/pricing/details/machine-learning/\)](https://azure.microsoft.com/en-us/pricing/details/machine-learning/).

- NVIDIA_TESLA_K80
- NVIDIA_TESLA_P100
- NVIDIA_TESLA_V100
- NVIDIA_TESLA_M60
- NVIDIA_TESLA_P40
- NVIDIA_TESLA_T4
- NVIDIA_TESLA_A100

3.

Google

- Ai platform notebook
- container registry
- kubeflow pipelines
- cloud build
- ai platform training
- ai platform prediction
- dataflow
- cloud storage
- cloud sql
- bigquery

Microsoft Azure ML

- Azure ML Pipelines
- Azure Container Instance
- Azure Kubernetes Service
- Azure ML datasets
- Interpretability
- Azure ML model registry
- Azure ML run
- Azure ML Supports alerts
- Azure ML provides monitoring
- Azure Data Factory

Amazon Sagemaker streamlines MLOps

- Pipelines supported
- Model Registry
- Projects - enables CI/CD across end to end lifecycle
- Model Monitor

IBM Cloud Pak

- WS watson studio
- Watson Machine Learning
- Watson Open Scale
- Watson Knowledge Catalog
- Data Virtualization
- CI/CD

4.

Amazon Sagemaker

- amazon sagemaker uses amazon cloudwatch which records application logs for 15 months

Google Cloud Vertex

- Cloud Logging service is used for application logging and resource logging

Microsoft Azure ML

- provides Azure Monitor Logs
- MLFlow can be used for logging

IBM

- platform logs which are used for logging app logs and resource logs

5.

Google

- AI Platform Training Jobs provides method to monitor training jobs in progress for metrics like accuracy and throughput

Amazon Sagemaker

- Amazon Cloud Watch can be utilized in order to monitor training jobs
- Sagemaker Console can also be used for monitoring training jobs

Microsoft Azure ML

- Azure ML provides abilities to monitor runs [link \(https://docs.microsoft.com/en-us/azure/machine-learning/how-to-track-monitor-analyze-runs?tabs=python#monitor-run-performance\)](https://docs.microsoft.com/en-us/azure/machine-learning/how-to-track-monitor-analyze-runs?tabs=python#monitor-run-performance).

IBM Cloud Pak (https://dataplatform.cloud.ibm.com/docs/content/wsj/analyze-data/ml_dlaas_working_with_new_models.html)

- Cloud Pak allows for monitoring of a training run by creating a training manifest file

6. Google AI

- autoscaling can be done through google compute engine clusters
- autoscaling can also be done for inference prediction as well through the google cloud ai platform

Amazon Sagemaker [link \(https://docs.aws.amazon.com/sagemaker/latest/dg/endpoint-auto-scaling.html\)](https://docs.aws.amazon.com/sagemaker/latest/dg/endpoint-auto-scaling.html).

- supports autoscaling in response to changes in workload, a scaling policy can be defined according to different metrics

Microsoft Azure ML [link \(https://docs.microsoft.com/en-us/azure/machine-learning/how-to-autoscale-endpoints?tabs=azure-cli\)](https://docs.microsoft.com/en-us/azure/machine-learning/how-to-autoscale-endpoints?tabs=azure-cli).

- Autoscale is supported which allows for autoscaling of a model which allows for setting various rules for autoscaling

IBM [link \(https://www.ibm.com/docs/en/cloud-paks/cp-data/4.0?topic=cluster-scaling-services\)](https://www.ibm.com/docs/en/cloud-paks/cp-data/4.0?topic=cluster-scaling-services).

- using the scaleConfig setting in IBM Cloud Pak

7.

In []: *#Google yaml file*

```
trainingInput:
  jobId: 'dl-test'
  scaleTier: CUSTOM
  masterType: complex_model_m
  workerType: complex_model_m
  parameterServerType: large_model
  workerCount: 9
  parameterServerCount: 3
  runtimeVersion: '2.6'
  packageUri: 'package_dir'
  pythonVersion: '3.7'
  region:
  scheduling:
    maxWaitTime: 3600s
    maxRunningTime: 7200s
```

```
In [ ]: #Sagemaker yaml file
apiVersion: sagemaker.aws.amazon.com/v1
kind: TrainingJob
metadata:
  name: 'dl-test'
spec:
  hyperParameters:
    - name: max_depth
      value: "5"
  algorithmSpecification:
    trainingImage: 'some image'
    trainingInputMode: File
  roleArn: arn:aws:iam::123456789012:role/service-role/AmazonSageMaker-ExecutionRole
  region: us-west-2
  outputDataConfig:
    s3OutputPath: ""
  resourceConfig:
    instanceCount: 1
    instanceType: ml.m4.xlarge
    volumeSizeInGB: 5
  stoppingCondition:
    maxRuntimeInSeconds: 7200
  inputDataConfig:
    - channelName: train
      dataSource:
        s3DataSource:
          s3DataType: S3Prefix
          s3Uri: ''
          s3DataDistributionType: FullyReplicated
        contentType: text/csv
        compressionType: None
    - channelName: validation
      dataSource:
        s3DataSource:
          s3DataType: S3Prefix
          s3Uri: ""
          s3DataDistributionType: FullyReplicated
        contentType: text/csv
        compressionType: None
  tags:
    - key: tagKey
      value: tagValue
```



```
In [ ]: #Azure ML
code:
  local_path: src
command: >-
  python main.py
  --dataset ${{inputs.iris_csv}}
  --C ${{inputs.C}}
  --kernel ${{inputs.kernel}}
  --coef0 ${{inputs.coef0}}
inputs:
  "some dataset"
  C: 0.8
  kernel: "rbf"
  coef0: 0.1
environment: azureml:AzureML-pytorch-0.24-ubuntu18.04-py37-cpu:9
compute: azureml:cpu-cluster
display_name: dl-test
experiment_name: dltest
description: Train a deep learning model
```

<https://raw.githubusercontent.com/aws/amazon-sagemaker-operator-for-k8s/master/samples/xgboost-mnist-trainingjob.yaml> (<https://raw.githubusercontent.com/aws/amazon-sagemaker-operator-for-k8s/master/samples/xgboost-mnist-trainingjob.yaml>)

```

In [ ]: # IBM Dataplatform
model_definition:
  framework:
    name: pytorch
    version: "1.1"
    runtimes:
      name: python
      version: "3.7"
  name: tf-mnist
  description: Simple MNIST model implemented in pytorch
  execution:
    command: python3 convolutional_network.py --trainImagesFile ${DATA_DIR}/train-images-idx3-ubyte.gz
      --trainLabelsFile ${DATA_DIR}/train-labels-idx1-ubyte.gz --test
ImagesFile ${DATA_DIR}/t10k-images-idx3-ubyte.gz
      --testLabelsFile ${DATA_DIR}/t10k-labels-idx1-ubyte.gz --learnin
ngRate 0.001
      --trainingIters 20000000
    compute_configuration:
      name: v100
  training_data_reference:
    name: MNIST image data files
    connection:
      endpoint_url: <auth-url>
      access_key_id: <username>
      secret_access_key: <password>
    source:
      bucket: mnist-training-data
      type: s3
  training_results_reference:
    name: DL Model Storage
    connection:
      endpoint_url: <auth-url>
      access_key_id: <username>
      secret_access_key: <password>
    target:
      bucket: mnist-training-models
      type: s3

```

There are common fields specifying python version, region, dataset used, hyperparameters, compute instance specifications, and some of this is obfuscated within an app image in GCP.

Problem 3

Section 1

a

✔ Create Project: jd3720-6998

Just now

[SELECT PROJECT](#)

b

Deployment name *

minikf-jd3720

Zone

northamerica-northeast1-a

?

Machine type

Machine family

GENERAL-PURPOSE

COMPUTE-OPTIMIZED

Machine types for common workloads, optimized for cost and flexibility


Series

N1

Powered by Intel Skylake CPU platform or one of its predecessors

Machine type

n1-standard-8 (8 vCPU, 30 GB memory)

	vCPU	Memory
	8	30 GB

✓ CPU PLATFORM AND GPU

Boot Disk

Boot disk type *

Standard Persistent Disk

?

Boot disk size in GB *

200

?

←

minikf-jd3720

⋮

✓ minikf-jd3720 has been deployed

Overview - minikf-jd3720

minikf minikf.jinja

minikf-jd3720-firewall firewall

minikf-jd3720-firewall-http01 firewall

minikf vm_instance.py

minikf-jd3720 vm instance

minikf-jd3720-data-disk disk

minikf-jd3720-password password.py

✕ minikf

MiniKF

Solution provided by Arrikto Inc.

MiniKF dashboard

<https://minikf-jd3720.endpoint.s.jd3720-6998.cloud.google/>

MiniKF username

user

MiniKF password

XS5884mwB7

Instance

[minikf-jd3720](#)

Instance zone

northamerica-northeast1-a

Instance machine type


n1-standard-8

✓ MORE ABOUT THE SOFTWARE

A screenshot of a Google Cloud SSH terminal window. The browser address bar at the top shows the URL: ssh.cloud.google.com/projects/jd3720-6998/zones/northamerica-northeast1-a/instances/minikf-jd3720?authuser=0&hl=en_US&proj... The terminal background is blue. A white box with a black border contains the following text:
Minikube + KubeFlow + Rok = MiniKF
Provisioning completed.
Log in to MiniKF here:
https://minikf-jd3720.endpoints.jd3720-6998.cloud.goog
Credentials:
* Username: user
* Password: XS5884mwB7



C

https://htmtopdf.herokuapp.com/ipynbviewer/temp/883e3a609a9dedfc7023cddb7db5ea57/HW5_6998.html?t=1638763941282 19/37








Section 2

a

 kubeflow-user (Owner) 

Notebook Servers

[+ NEW SERVER](#)

Status	Name	Type	Age	Image	GPUs	CPUs	Memory	Volumes	
	tutorialjd3720		just now	tensorflow-1.14.0-notebook-cpu:kubecon-work...	0	0.5	1Gi		CONNECT  

b

```
groups: cannot find name for group ID 1337
jovyan@tutorialjd3720-0:~$ cd data
jovyan@tutorialjd3720-0:~/data$ git clone -b kubecon-workshop https://github.com/kubeflow-kale/examples
Cloning into 'examples'...
remote: Enumerating objects: 178, done.
remote: Counting objects: 100% (9/9), done.
remote: Compressing objects: 100% (5/5), done.
remote: Total 178 (delta 5), reused 4 (delta 4), pack-reused 169
Receiving objects: 100% (178/178), 927.77 KiB | 10.91 MiB/s, done.
Resolving deltas: 100% (82/82), done.
jovyan@tutorialjd3720-0:~/data$
```

c

Kale Deployment Panel

Enable ☒

Pipeline Metadata

Select experiment

+ New Experiment

Experiment Name

jd3720

Pipeline Name

cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a

Run

HP Tuning with Katib ☐

SET UP KATIB JOB

Advanced Settings

COMPILE AND RUN

Terminal 1

cifar10_classification.ipynb

Code

step: testwhole depends on: ●

```
[19]: correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data,
        total += labels.size(0)
        correct += (predicted == labels).sum())

print('Accuracy of the network on the 10000 te
100 * correct / total))
```


Accuracy of the network on the 10000 test imag

```
[20]: class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1


for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / c
```

```
Accuracy of plane : 49 %
Accuracy of car : 55 %
Accuracy of bird : 34 %
Accuracy of cat : 28 %
Accuracy of deer : 35 %
Accuracy of dog : 33 %
Accuracy of frog : 84 %
Accuracy of horse : 57 %
Accuracy of ship : 79 %
Accuracy of truck : 66 %
```


[]:





File Edit View Run Kernel Tabs Settings Help



Kale Deployment Panel




Enable 



Pipeline Metadata

Select experiment

+ New Experiment 

Experiment Name


jd3720

Pipeline Name


cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a





Run



HP Tuning with Katib 


SET UP KATIB JOB


Advanced Settings










Validating notebook... 


Taking snapshot... 



Terminal 1 

cifar10_classification.ipynb 



Code 

Accuracy of deer : 50 %
Accuracy of dog : 50 %
Accuracy of frog : 66 %
Accuracy of horse : 61 %
Accuracy of ship : 65 %
Accuracy of truck : 58 %

[]:

https://htmtopdf.herokuapp.com/ipynbviewer/temp/883e3a609a9dedfc7023cddb7db5ea57/HW5_6998.html?t=1638763941282

23/37

Kale Deployment Panel

Enable ☒

Pipeline Metadata

Select experiment

+ New Experiment

Experiment Name

jd3720

Pipeline Name

cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a

Run

HP Tuning with Katib ☐

SET UP KATIB JOB

Advanced Settings

Validating notebook... ☒

Taking snapshot... [Done](#)

Compiling notebook... [Done](#) ☒

Uploading pipeline... [Done](#)

Running pipeline... [View](#)

COMPILE AND RUN

Terminal 1

cifar10_classification.ipynb

Code

Python 3

step: train depends on: ●

```
[12]: import torch.optim as optim

net = Net()
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum
```

Train

```
[13]: for epoch in range(TRAIN_STEPS): # loop over the dataset

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

print('Finished Training')
```

```
[1, 2000] loss: 2.230
[1, 4000] loss: 1.892
[1, 6000] loss: 1.677
[1, 8000] loss: 1.567
[1, 10000] loss: 1.521
[1, 12000] loss: 1.459
[2, 2000] loss: 1.400
[2, 4000] loss: 1.383
[2, 6000] loss: 1.356
[2, 8000] loss: 1.315
[2, 10000] loss: 1.331
```

Kale Deployment Panel

Enable ☒

Pipeline Metadata

Select experiment

+ New Experiment

Experiment Name

jd3720

Pipeline Name

cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a

Run

HP Tuning with Katib ☐

SET UP KATIB JOB

Advanced Settings

Validating notebook... ✓

Taking snapshot... [Done](#) [🔗](#)

Compiling notebook... [Done](#) ✓

Terminal 1

cifar10_classification.ipynb

Python 3

```
[1, 4000] loss: 1.892
[1, 6000] loss: 1.677
[1, 8000] loss: 1.567
[1, 10000] loss: 1.521
[1, 12000] loss: 1.459
[2, 2000] loss: 1.400
[2, 4000] loss: 1.383
[2, 6000] loss: 1.356
[2, 8000] loss: 1.315
[2, 10000] loss: 1.331
[2, 12000] loss: 1.293
Finished Training
```

Test on test data

step: **testontest** depends on: ●

```
[14]: dataiter = iter(testloader)
      images, labels = dataiter.next()

      # print images
      imshow(torchvision.utils.make_grid(images))
      print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



```
GroundTruth:  cat  ship  ship plane
```

```
[15]: outputs = net(images)

[*]: _, predicted = torch.max(outputs, 1)

      print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                                     for j in range(4)))
```

```
Predicted:  cat  ship  ship plane
```

Kale Deployment Panel

Enable ☒

Pipeline Metadata

Select experiment

+ New Experiment

Experiment Name

jd3720

Pipeline Name

cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a

Run

HP Tuning with Katib ☐

SET UP KATIB JOB

Advanced Settings

Validating notebook... ☒

Taking snapshot... [Done](#) ☒

Compiling notebook... [Done](#) ☒

Uploading pipeline... [Done](#) ☒

Running pipeline... [View](#) ☒

COMPILE AND RUN

Terminal 1

cifar10_classification.ipynb

Python 3

Predicted: cat snip snip plane

Performance on whole dataset

step: testwhole depends on: ●

```
[*]: correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d
100 * correct / total))
```

Accuracy of the network on the 10000 test images: 54 %

```
[*]: class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]
```

Accuracy of plane : 66 %
 Accuracy of car : 62 %
 Accuracy of bird : 42 %
 Accuracy of cat : 23 %
 Accuracy of deer : 50 %
 Accuracy of dog : 50 %
 Accuracy of frog : 66 %
 Accuracy of horse : 61 %
 Accuracy of ship : 65 %

Kale Deployment Panel

Enable



Pipeline Metadata

Select experiment

+ New Experiment

Experiment Name

jd3720

Pipeline Name

cifar10-classification

Pipeline Description

Sequential PyTorch pipeline to train a

Run

HP Tuning with Katib



SET UP KATIB JOB

Advanced Settings

Validating notebook...

Taking snapshot... [Done](#)

Compiling notebook... [Done](#)

Terminal 1

cifar10_classification.ipynb

Python

Downloading on.tar.gz visual 3/3

100.0%

Extracting ./data/cifar-10-python.tar.gz to ./data

Files already downloaded and verified

Visualize dataset

functions

[9]:

```
import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

skip

[10]:

```
# get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(
```

bird plane ship horse

The screenshot shows the Kale Deployment Panel on the left and a Jupyter Notebook on the right.

Kale Deployment Panel:

- Enable:** Toggle switch is turned on.
- Pipeline Metadata:**
 - Select experiment: **jd3720**
 - + New Experiment
 - Experiment Name: **jd3720**
 - Pipeline Name: **cifar10-classification**
 - Pipeline Description: **Sequential PyTorch pipeline to train a**
- Run:**
 - HP Tuning with Katib: Toggle switch is turned off.
 - SET UP KATIB JOB
- Advanced Settings**

Jupyter Notebook (cifar10_classification.ipynb):

```
[7]: TRAIN_STEPS = 2
```

Load and transform dataset

step: **dataprocessing**

```
[8]: input_data_folder = "./data"

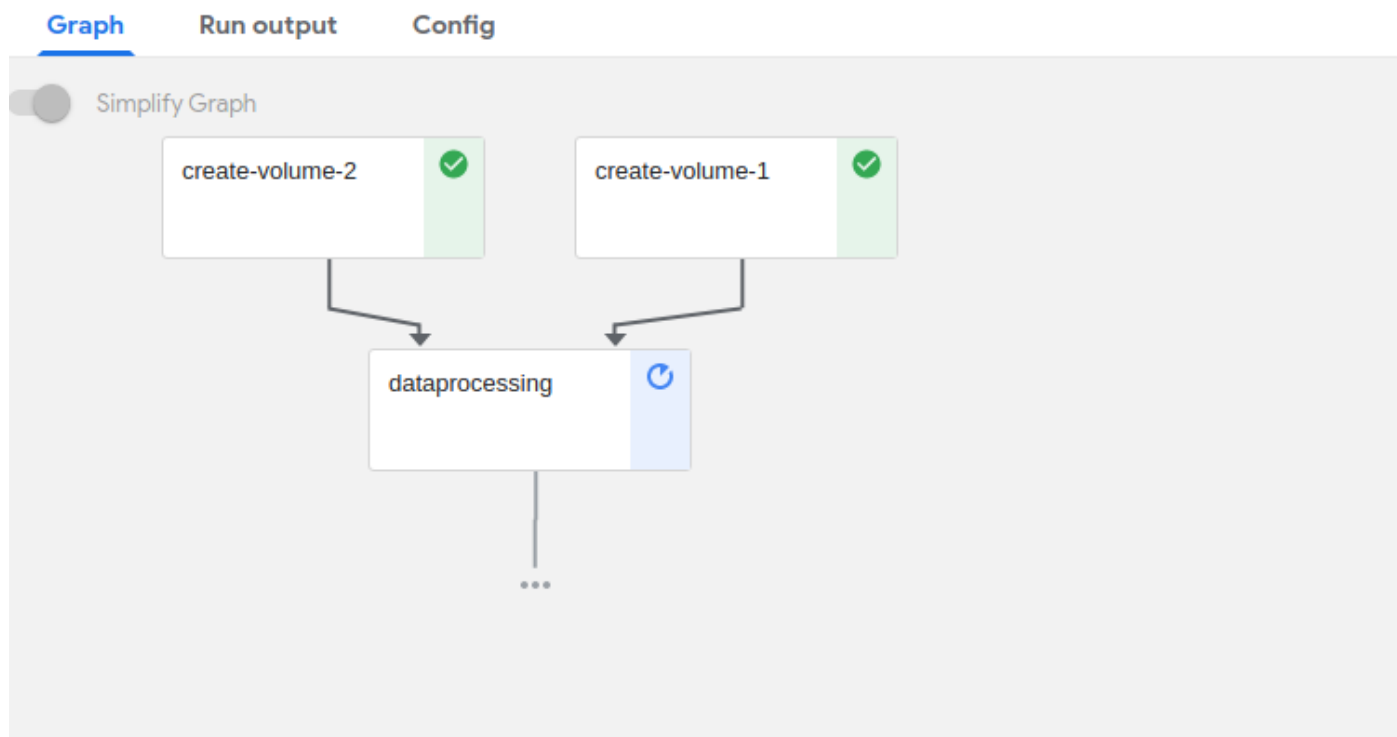
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

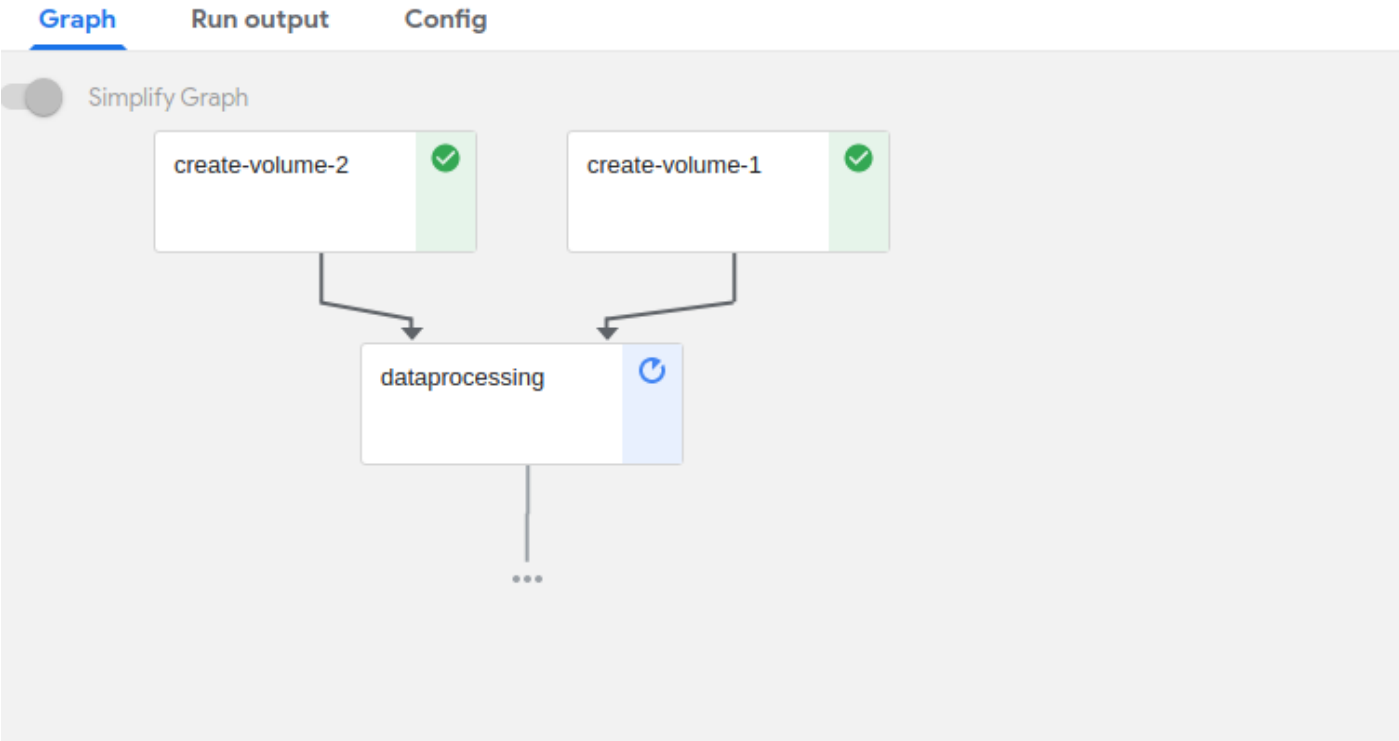
trainset = torchvision.datasets.CIFAR10(root=input_data_folder,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=64,
                                           shuffle=True, num_workers=4)

testset = torchvision.datasets.CIFAR10(root=input_data_folder,
                                       download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=64,
                                         shuffle=False, num_workers=4)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz
100.0%
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

Visualize dataset





d

Experiments > jd3720

← ✓ cifar10-classification-rzhe0-1r3lg

Graph

Run output

Config

☒ Simplify Graph

create-volume-2 ✓

create-volume-1 ✓

dataprocessing ✓

train ✓

testontest ✓

testwhole ✓

testwhole



Rok autosnapshot

Rok has successfully created a snapshot for step testwhole.

To **explore the execution state** at the **beginning** of this step follow the instructions below:

1. View the [snapshot in the Rok UI](#).
2. Copy the Rok URL.
3. Create a new Notebook Server by using this Rok URL to autofill the form.

testwhole



```
2021-12-03 20:22:00 Kale marshalling [INF
```

```
2021-12-03 20:22:05 Kale marshalling [INF
```

```
Accuracy of the network on the 10000 test images:
```

```
Accuracy of plane : 78 %  
Accuracy of car : 63 %  
Accuracy of bird : 59 %  
Accuracy of cat : 26 %  
Accuracy of deer : 48 %  
Accuracy of dog : 37 %  
Accuracy of frog : 63 %  
Accuracy of horse : 43 %  
Accuracy of ship : 56 %  
Accuracy of truck : 36 %
```

dataprocessing



Rok autosnapshot

Rok has successfully created a snapshot for step `dataprocessing`.

To **explore the execution state** at the **beginning** of this step follow the instructions below:

1. View the [snapshot in the Rok UI](#).
2. Copy the Rok URL.
3. Create a new Notebook Server by using this Rok URL to autofill the form.

dataprocessing



```
Files already downloaded and verified
```

```
Files already downloaded and verified
```

```
2021-12-03 20:16:22 Kale marshalling [INFO]
```

```
2021-12-03 20:16:22 Kale marshalling [INFO]
```


Rok autosnapshot

Rok has successfully created a snapshot for step train.

To **explore the execution state** at the **beginning** of this step follow the instructions below:

1. View the [snapshot in the Rok UI](#).
2. Copy the Rok URL.
3. Create a new Notebook Server by using this Rok URL to autofill the form.

train

```
[1, 6000] loss: 1.633
[1, 8000] loss: 1.530
[1, 10000] loss: 1.487
[1, 12000] loss: 1.463
[2, 2000] loss: 1.379
[2, 4000] loss: 1.360
[2, 6000] loss: 1.351
[2, 8000] loss: 1.331
[2, 10000] loss: 1.304
[2, 12000] loss: 1.264
```

Finished Training

2021-12-03 20:20:11 Kale marshalling

testontest



```
2021-12-03 20:21:01 Kale marshalling [INF
2021-12-03 20:21:08 Kale marshalling [INF
```



GroundTruth: cat ship ship plane

Predicted: bird ship ship plane

```
2021-12-03 20:21:10 Kale marshalling [INF
2021-12-03 20:21:10 Kale marshalling [INF
```

e

Deployments

+

DEPLOY MARKETPLACE SOLUTION

■ STOP

🗑

DELETE

📘

[Learn how to convert your deployment to KRM or Terraform](#)

☰ Filter

Filter deployments

<div>✓</div>	<div>●</div>	Name ↑	Created on	Last modified	Labels
<div>✓</div>	<div>✓</div>	minikf-jd3720	12 hours ago	12 hours ago	cloud-mark... : arrikto-pu... <div>▼</div>

Delete deployment

🎯

Delete minikf-jd3720 and all resources created by it, such as VMs, load balancers and disks

○

Delete minikf-jd3720, but keep resources created by it

Problem 4

1.

Episodic tasks are those that have an end state or terminal state. An episodic task can be a game like checkers which has a clear end for an agent. Continuous task is when there is no end state and the task doesn't end. An example of a continuous task is a surveillance task where the agent must constantly monitor and identify intruders or changes to the environment.

2.

Exploration in RL means taking an action that has not been taken before while exploitation is taking an action which leads to the greatest rewards (best action) given the learner's knowledge at this point in time.

Epsilon greedy aids in keeping the network from getting stuck in sub-optimal policies while at the same time allowing the network to build off of the best policy found at this point in time. As such epsilon greedy aids in exploring the state space.

Epsilon should follow a schedule during deep RL training as the network doesn't have any knowledge of the best course of action and exploitation can't be done. As epsilon should be high in the beginning and become smaller. In the beginning exploring is almost always done and then once learnings have taken place exploitation can occur given the network has now learned. As such a small epsilon is needed for enabling some exploration.

3.

Initialize replay memory D to capacity N

- init replay memory

Initialize action-value function Q with random weights

- init Q function

for episode = 1, M do:

- loop over episodes

Initialise sequence $s_1 = x_1$ and preprocessed sequence $\varphi_1 = \varphi(s_1)$

- init start state and history

for $t = 1, T$ do:

- start until terminal state is reached

With probability epsilon select a random action a_t otherwise select $a_t = \arg \max_a Q(s_t, a; \theta)$

- use an epsilon greedy policy where the exploit/exploitation takes place

Execute action a_t in emulator and observe reward r_t and image $x(t+1)$:

- action and take into account reward and the next state

Set $s(t+1) = s_t, a_t, x(t+1)$ and preprocess $\varphi(t+1) = \varphi(s(t+1))$

- update state

Store transition $(\varphi_t, a_t, r_t, \varphi_t + 1)$ in D

- store the new experience

Sample random minibatch of transitions $(\varphi_j, a_j, r_j, \varphi(j + 1))$ from D

- sample a minibatch of experiences from the replay memory

Set $y_j = \{r_j + \lambda \max Q\}$:

- set the target value equal to the reward of the state s_j for terminal states or set the target value equal to the reward + future discounted rewards for non terminal states

perform a gradient descent step

- finally update the Q network

4

The target Q networks keeps the target distribution stable so that the learner is not always trying to fit to a moving distribution.

5

Deep learning algorithms expect data to be independent but in online Reinforcement Learning it is that the data is interdependent because future steps are dependent on previous steps. As such experience replay stores previous steps and the learner is able to randomly sample from this distribution which solves the problem of correlated data and non-stationary distributions

6

Instead of drawing experiences at random from memory, experience which carry the most information and provide the learned with the best signal are chosen.

7

Similar

- various actors
- shared replay memory
- each respective actor have their own instance of the environment

Differences

- ape-x uses prioritized experience replay, Gorilla doesn't
- gorilla uses parameter server and ape-x doesn't
- gorilla has multiple learners on multiple GPUs while apex has a single learner on a GPU